

CAS CS 561: Data Systems Architectures

Boston University

Spring 2021

Class Syllabus

Course Description: Data is everywhere. As scientists, users, and citizens we are both generating and exploiting large, ever-growing, diverse sets of data. For several applications – ranging from scientific discovery to business analysis, governance, and every-day activities – we are directly using and indirectly affecting hundreds of data systems! The big challenge is to turn data into useful knowledge, and to do so quickly, in order to increase the impact of the new insights. Achieving these goals comes with a number of technical challenges. How to exploit the continuously evolving hardware (storage, computation, network)? How to collect all incoming data efficiently? How to query dynamic collections of data that keep accumulating incoming data? How to parallelize query processing from one core to a few (scale-up), and then to thousands (scale out)? What are the needs of evolving workloads (hybrid transactional/analytical processing, graph analytics, Internet-of-Things, micro-payments, monitoring)? In this course, we will discuss how to design data systems that can address these challenges. We will see in detail the two driving forces behind innovation in data systems: hardware and workloads, and we will discuss recent and future trends of both. We will use examples from several data management areas including relational systems, distributed database systems, key value stores, newSQL and NoSQL systems, data systems for machine learning (and machine learning for data systems), interactive analytics, and data management as a service. In a quickly moving industry and research landscape, such skills are essential.

Prerequisites: The class requires familiarity with database systems at the level **CS460/660**, and with algorithms, data structures, computer systems, and system programming at the level of **CS210**. Please see the instructor if you are not sure about the level of your preparation.

Instructor: Manos Athanassoulis (mathan@bu.edu)

office hours: TBA

office: MCS 106

Teaching Assistants: Andy Huynh (ndhuynh@bu.edu)

Meeting Times and Places

lectures: Tu/Th, 12:30-1:45 pm, CAS 116

lab: F, 9:05-9:55 am, CAS 213

Course Website: <https://bu-disc.github.io/CS561/>

Textbooks (not required): There is no textbook that covers cutting edge research, however, the data management community has produced top quality textbooks that can serve as reference to provide background material. The class is based on recent research papers which will be available to you through the BU network.

- R. Ramakrishnan and J. Gehrke. [Database Management Systems](#). Third Edition. McGraw-Hill 2002.

An excellent collection of classic papers in the database field is the following:

- [Readings in Database Systems](#). P. Bailis, J. Hellerstein, M. Stonebraker, editors.

Other good background material is:

- Architecture of a Database System. By J. Hellerstein, M. Stonebraker and J. Hamilton. Foundations and Trends in Databases, 2007
- The Design and Implementation of Modern Column-store Database Systems. By D. Abadi, P. Boncz, S. Harizopoulos, S. Idreos, S. Madden. Foundations and Trends in Databases, 2013
- Massively Parallel Databases and MapReduce Systems. By Shivnath Babu and Herodotos Herodotou. Foundations and Trends in Databases, 2013

Course Requirements

Paper Presentation: After the initial 4 classes (2 weeks) all students will take turns presenting papers. In each class we will discuss one or two main paper(s) (and there will be a few background papers), and each student will present once in the semester, either alone or as a group of two students). The student(s) presenting will be responsible to outline strong and weak points of the paper and propose at least one idea for improving the approach presented in the paper. **All students** will read the presented paper.

Paper Reviews: All students should read all papers. *Reading the paper and writing a review is very important to help the students prepare for the class presentation and discussion.* Every student is expected to deliver a **review for 8 papers** and **answer a technical question for all other papers**. Each paper will be clearly marked as a paper for review or a paper for a technical question (which will be clearly asked). Throughout the semester each student can skip 2 reviews (which means 10 papers will be marked for review in total). Every review/answer for a given paper has to be handed-in **before the class**, having the class starting time as a hard deadline.

A **review** consists of a few paragraphs answering the following questions: (i) what is the problem, (ii) why it is important, (iii) why is it hard, (iv) why older approaches are not enough, (v) what is key idea and why it works (a list of at least three key points), (vi) what might be missing and how can we improve this idea (a list of at least three key points), (vi) an evaluation as to whether the paper supports its claims, and (vii) possible next steps of the work presented in the paper. The ideal size of the review is about 1 page, single column, 10pt font, 1 inch margin (and it may only exceed 1 page if the student wants to elaborate on how to improve the ideas on the paper). Every paper should have at least two reviewers.

Project: Finally, this class requires a semester-long project and a final report in the style of a conference paper. The project will be either implementation-heavy or research-oriented. Students will work in teams of 2 (implementation project) or 3 (research project) and after the first two weeks each team will have been associated with a specific project. Students can propose their own

research project upon approval by the instructor.

The overall grade will be based on the following policy:

- Class participation: 5%
- Paper reviews: 20%
- Technical questions: 5%
- Paper presentation: 25%
- Project proposal: 5%
- Mid-semester project progress report: 5%
- Project: 35%

Topics: Throughout the class we will cover data systems design principles from the following different angles.

1. What affects new data systems designs (data and applications, emerging hardware, and new workloads)
2. Traditional Data Systems for Modern Hardware
3. Distributed Database Systems
4. Scale-out Systems: from Map-Reduce to SQL-on-Hadoop
5. NoSQL, NewSQL and Key-Value stores

Important Dates for all classes (more info at: <https://www.bu.edu/reg/calendars/semester/>)

February 5th, last day to add a class

March 1st, last day to drop (without a “W”)

Important Dates for CS561

February 9th, last day to select a project

February 23rd, submit the early project report

April 30th, submit final project report

Tentative Schedule

Week 1	Lecture 1	Introduction to Data Systems and CS561
	Lecture 2	Data Systems Architectures Essentials – Part 1
Week 2	Lecture 3	Data Systems Architectures Essentials – Part 1
	Lecture 4	Class Project Overview
Week 3	Lecture 5	Storage Layouts: Row-Stores vs. Column-Stores
	Lecture 6	Storage Layouts: Adaptive & Hybrid Layouts
Week 4	Lecture 7	New Hardware: Data Systems for Flash & SMR Disks
	Lecture 8	New Hardware: Data Systems for Multi-Core
Week 5	No Class, Monday replacement day	
	Lecture 9	Indexing A: B+-Trees, Bitmaps, Hash-Index
Week 6	Lecture 10	Indexing B: Access Path Selection
	Lecture 11	Modern Storage Engines: HTAP Systems
Week 7	Lecture 12	Modern Storage Engines: Log-Structured Merge Trees
	Research talk 1	TBA (visiting lecture)
Week 8	No Class – Spring Break	

Week 9	Lecture 13	Indexing C: Data Skipping
	Lecture 14	Indexing D: Adaptive Indexing
Week 10	Research talk 2	TBA (visiting lecture)
	Research talk 3	TBA (visiting lecture)
Week 11	Lecture 15	In-Situ Data Processing: Efficiently Accessing Raw Data Files
	Lecture 16	Scientific Databases: Multi-dimensional Data Management
Week 12	Lecture 17	Distributed Databases: Global Distributed Systems
	Lecture 18	Map/Reduce: Data Management at Scale
Week 13	Lecture 19	Data Systems for ML: Data Processing Primitives for ML
	Lecture 20	ML for Data Systems: Automatic Data System Design
Week 14	Lecture 21	Indexing E: Learned Indexes
	Lecture 22	Indexing F: The Periodic Table of Access Methods
Week 15	Lecture 23	Project Presentations A
	Lecture 24	Project Presentations B

Collaboration Policy

You are strongly encouraged to collaborate with one another in studying the lecture materials and preparing for reviews and presentations.

You may discuss ideas and approaches to the projects with others (provided that you acknowledge doing so in your solution), but such discussions should be kept at a high level, and should not involve actual details of the code or of other types of answers. **You must complete the actual solutions on your own.**

Academic Misconduct

We will assume that you understand BU's Academic Conduct Code:

<http://www.bu.edu/academics/policies/academic-conduct-code>

Prohibited behaviors include:

- copying all or part of someone else's work, even if you subsequently modify it; this includes cases in which someone tells you what you should write for your solution
- viewing all or part of someone else's work
- showing all or part of your work to another student
- consulting solutions from past semesters, or those found online or in books
- posting your work where others can view it (e.g., online).

Incidents of academic misconduct will be reported to the Academic Conduct Committee (ACC). The ACC may suspend/expel students found guilty of misconduct. ***At a minimum, students who engage in misconduct will have their final grade reduced by one letter grade (e.g., from a B to a C).***