



PUMP UP THE VOLUME: PROCESSING LARGE DATA ON GPUS WITH FAST INTERCONNECTS

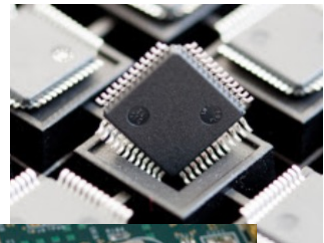
CLEMENS LUTZ ET.AL

Presented by : Syahril Dahler

BACKGROUND

Introduction of Co Processors

- FPGA (field-programmable gate array)
- ASIC (application-specific integrated circuit)
- GPU (Graphic Processing Unit)



GPU ARCHITECTURE

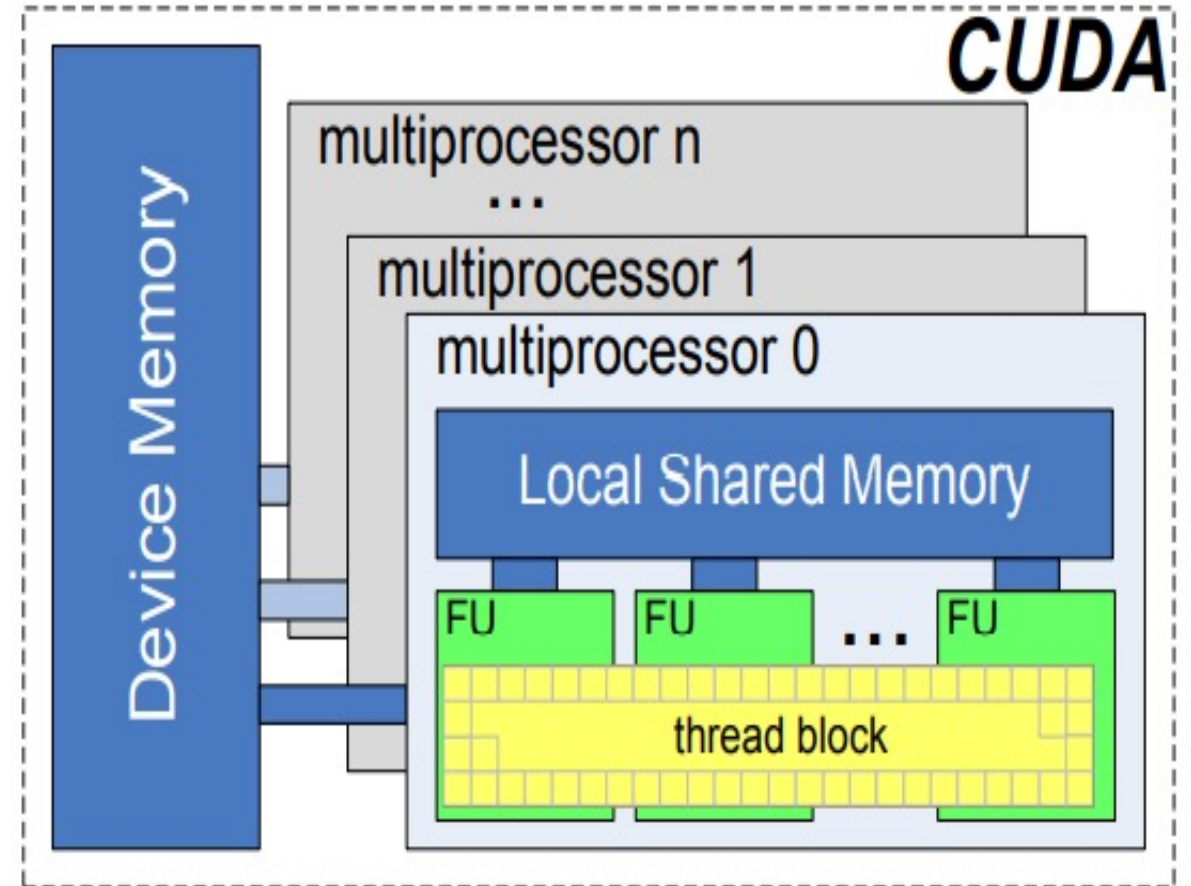
Tesla M2070 Processor:

Streaming Multiprocessors (SM): 14

Streaming Processors on each SM: 32

Total cores = $14 \times 32 = 448$ cores

Each Streaming Multiprocessor supports
1024 threads.



Compute unified device architecture

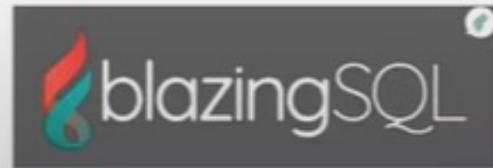
ADVANTAGES OF GPU

Parallel
Processing



Faster
result

GPU IN MODERN DATABASES



kinetica

SQREAM

CHALLENGES OF USING GPU FOR DATABASE APPLICATION

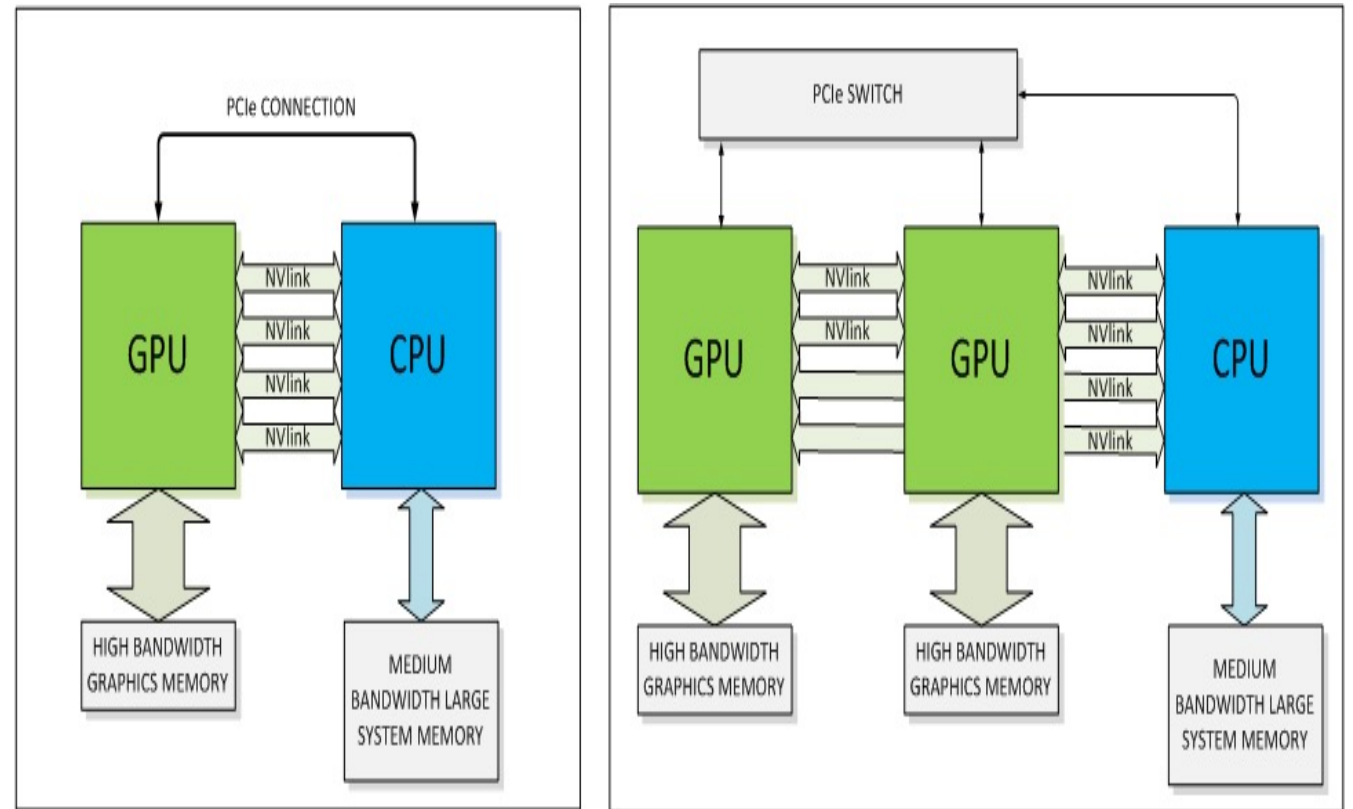
Transfer Bottleneck

- Low Interconnect Bandwidth
- Small GPU memory capacity
- Coarse grain cooperation of CPU and GPU
- **How to access data in main memory from GPU?**

FAST INTERCONNECT

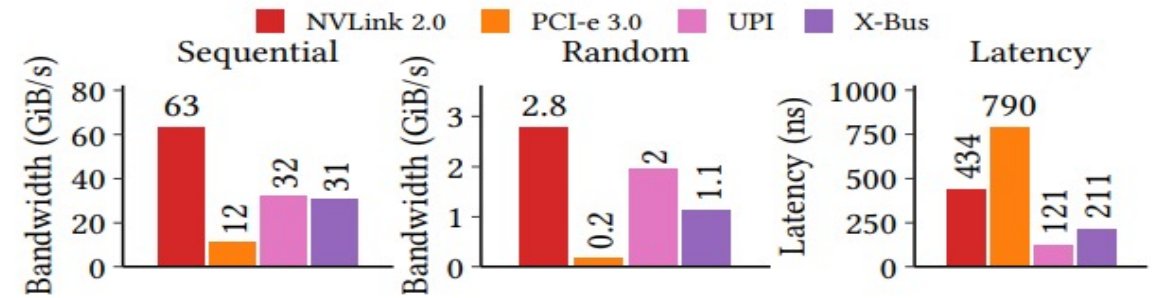
Faster interconnects
help to remedy transfer
bottleneck issues

NVLink 2.0

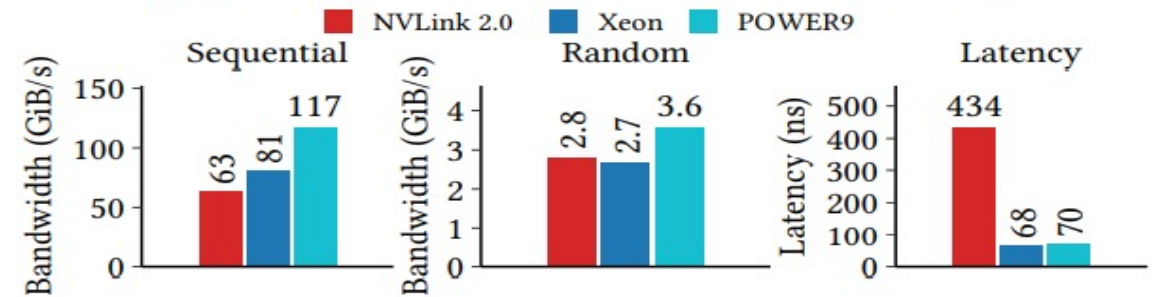


ANALYSIS OF A FAST INTERCONNECT

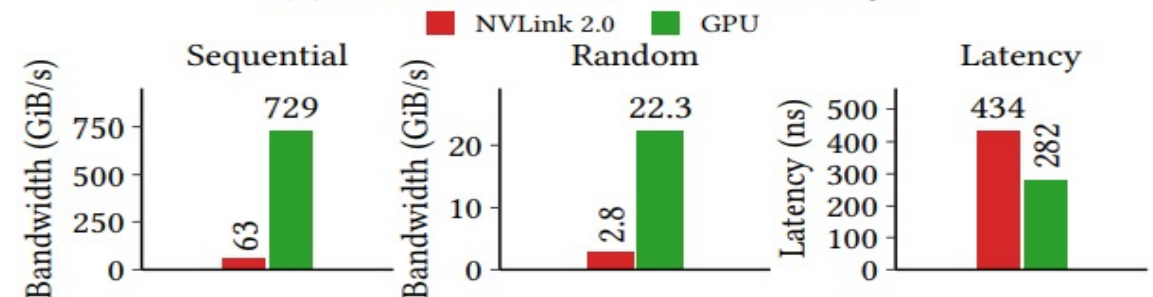
NVLink 2.0 improves the GPU's interconnect performance (data transfer)



(a) NVLink 2.0 vs. CPU & GPU Interconnects.



(b) NVLink 2.0 vs. CPU memory.



(c) NVLink 2.0 vs. GPU memory.

CHALLENGES DESPITE FAST CONNECTS (NV LINK) FOR QUERY PROCESSING

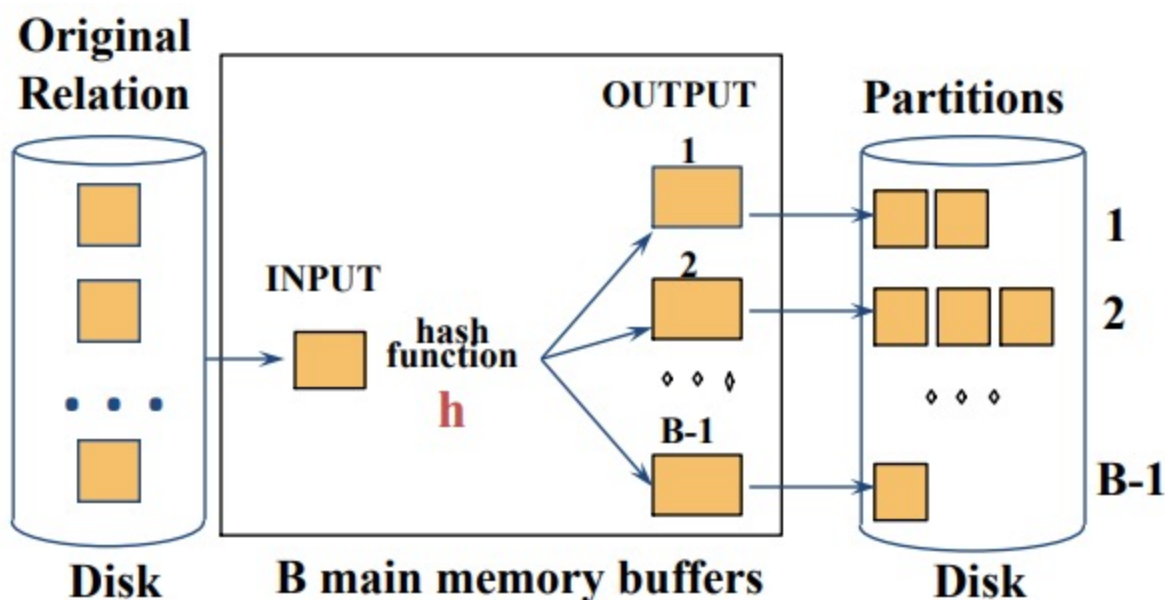
- ❑ Out-of-core GPU join operator must perform both data access and computation efficiently
- ❑ Join CPU and GPU requires effective cooperation. Locality and synchronization cost
- ❑ Increase build side → increase NP -HJ → spill Hash table to CPU memory → more irregular access to CPU memory (inefficient)

GOAL OF THE PAPER

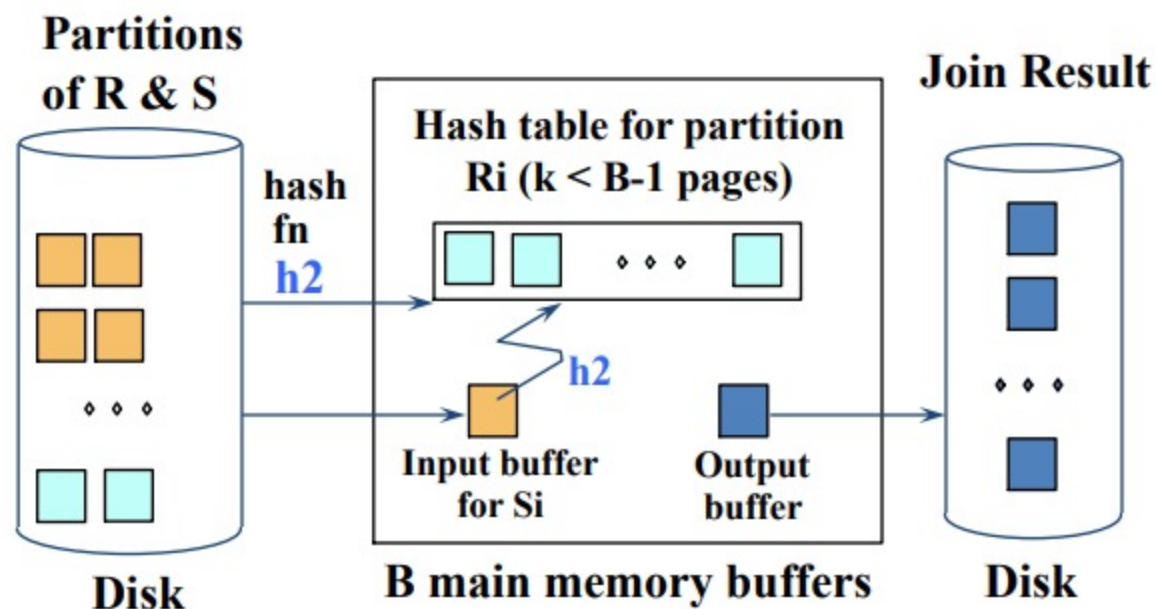
“Scale up GPU-accelerated data management to arbitrary data volumes”

Hash-Join

Partition both relations using hash function h : R tuples in partition i will **only** match S tuples in partition i



Read in a partition of R , hash it using h_2 ($\neq h$). Scan matching partition of S , probe hash table for matches



NO PARTITION HASH JOIN ($R \bowtie S$)

To take advantage of multi core processing

Build

Scan relation R and create a hash table on join key.

Probe

For each tuple in S, look up its join key in hash table for R. If a match is found, output combined tuple.

DATA TRANSFER BETWEEN CPU AND GPU

Push => CPU push
data to GPU

Pull => GPU pull
data from CPU

Table 1: An overview of GPU transfer methods.

| Method | Semantics | Level | Granularity | Memory |
|-----------------|-----------|-------|-------------|----------|
| Pageable Copy | Push | SW | Chunk | Pageable |
| Staged Copy | | | | |
| Dynamic Pinning | | | | |
| Pinned Copy | | | | Pinned |
| UM Prefetch | Pull | OS | Page | Unified |
| UM Migration | | | | Unified |
| Zero-Copy | | HW | Byte | Pinned |
| Coherence | | | | Pageable |

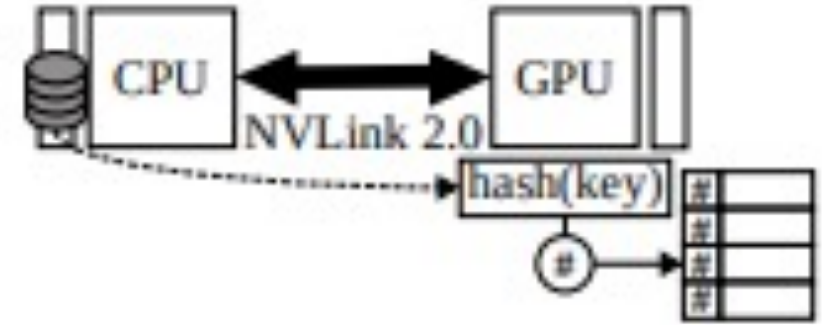
Coherence : GPU can directly access any CPU memory during execution (because of NVLink)

SCALING GPU HASH JOIN : SCALING PROBE SIZE

1. Build hash table on GPU by pulling R tuples on demand from CPU

2. Using Coherence transfer

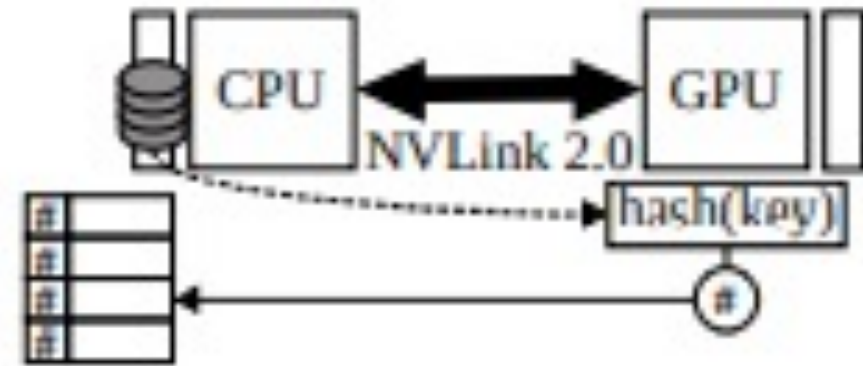
Baseline: data is copied into GPU memory to build hash table



(b) Data in CPU memory and hash table in GPU memory.

SCALING GPU HASH JOIN : SCALING BUILD SIZE

1. Hash Table is stored in CPU memory
2. No longer constrained by the GPU's memory capacity

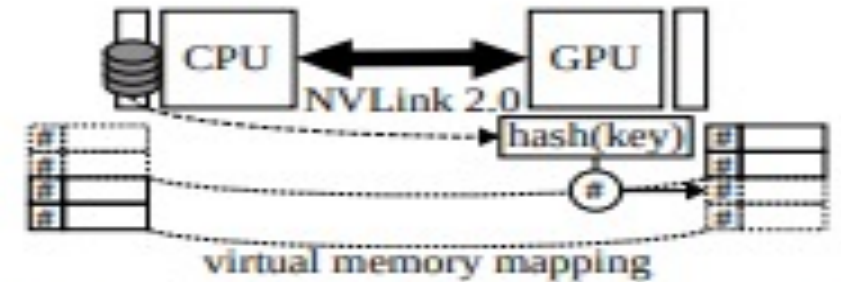


(a) Data and hash table in CPU memory.

SCALING GPU HASH JOIN : OPTIMIZE HASH TABLE PLACEMENT

Since GPU is much faster than CPU:

1. Place Hash Table on GPU memory and then spill to CPU memory
2. It is done by using Hybrid Hash table
3. Hybrid hash table uses virtual memory to abstract the physical location of memory page



(b) Data in CPU memory and hash table spills from GPU memory into CPU memory.

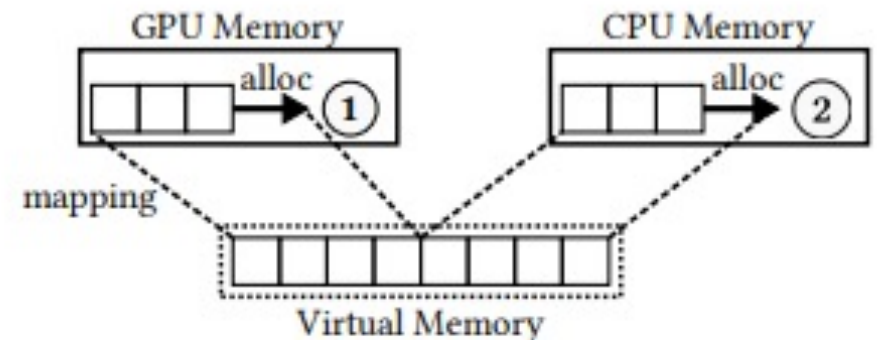


Figure 8: Allocating the hybrid hash table.

SCALING-UP USING CPU AND GPU : TASK SCHEDULING

To solve **load imbalance** issue

1. Adapt the CPU-oriented, **morsel-driven** approach
2. Give each processor the right amount of work to minimize execution skew by considering the increased latency of scheduling work on a GPU, and the higher processing rate of the GPU

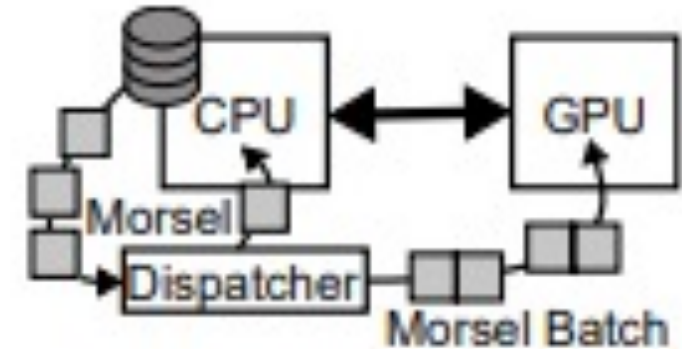
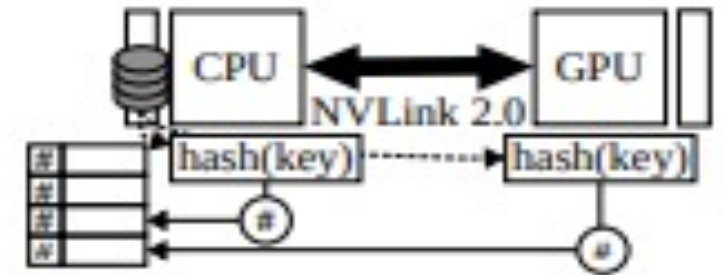


Figure 10: Dynamically scheduling tasks to CPU and GPU processors.

SCALING-UP USING CPU AND GPU : HETEROGENEOUS HASH TABLE PLACEMENT

A. CPU and GPU processing a join using a globally shared hash table (Het strategy) Same as scaling build size



(a) Cooperatively process join on CPU and GPU with hash table in CPU memory.

SCALING-UP USING CPU AND GPU : HETEROGENEOUS HASH TABLE PLACEMENT

Processors are fastest when accessing their local memories

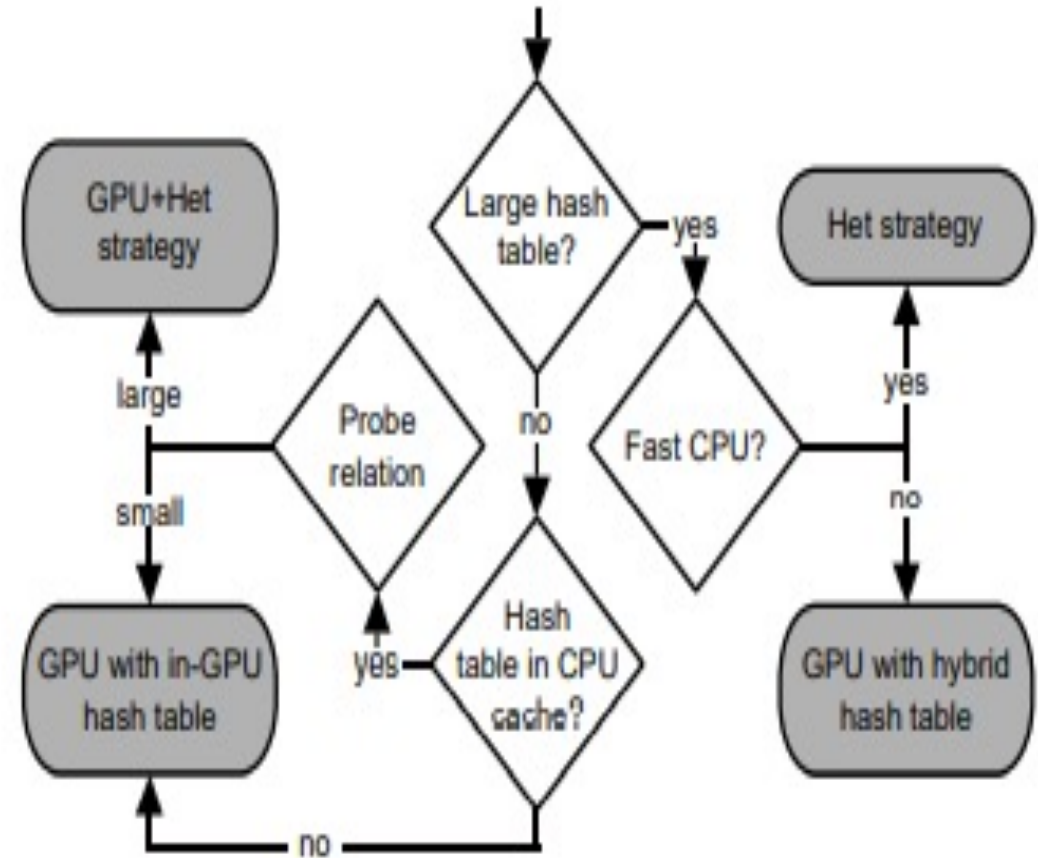
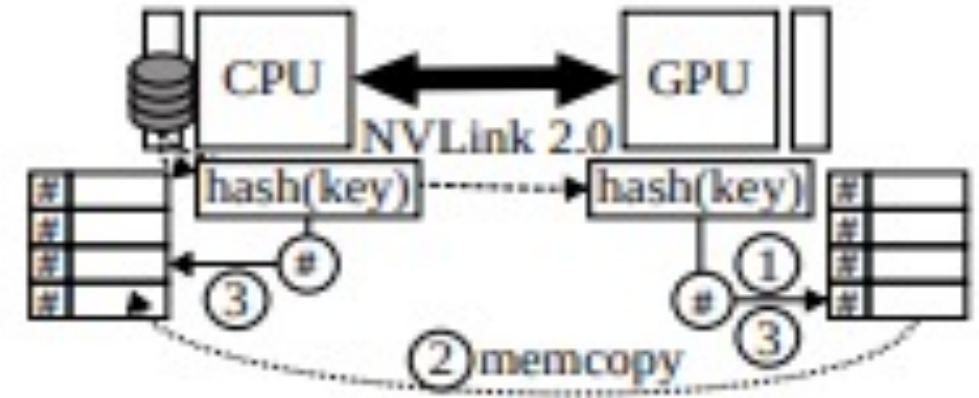


Figure 11: Hash table placement decision.

SCALING-UP USING CPU AND GPU : HETEROGENEOUS HASH TABLE PLACEMENT (WHEN HASH IS SMALL)

1. GPU build hash table in local memory
2. Copy to all other processors
3. Execute the probe phase on all processors using our heterogeneous scheduling strategy.



(b) Build hash table on GPU, copy the hash table to processor-local memories, and then cooperatively probe on CPU and GPU.

MULTI GPU HASH TABLE PLACEMENT

Advantages of multi-GPU

1. Using only GPUs avoids computational skew
2. Distributing large hash tables within GPU memory frees CPU memory bandwidth for loading the base relations
3. interleaving the hash table over multiple GPUs utilizes the full bi-directional bandwidth of fast interconnects, as opposed to the mostly uni-directional traffic of the Het strategy

EXPERIMENT: WORKLOADS

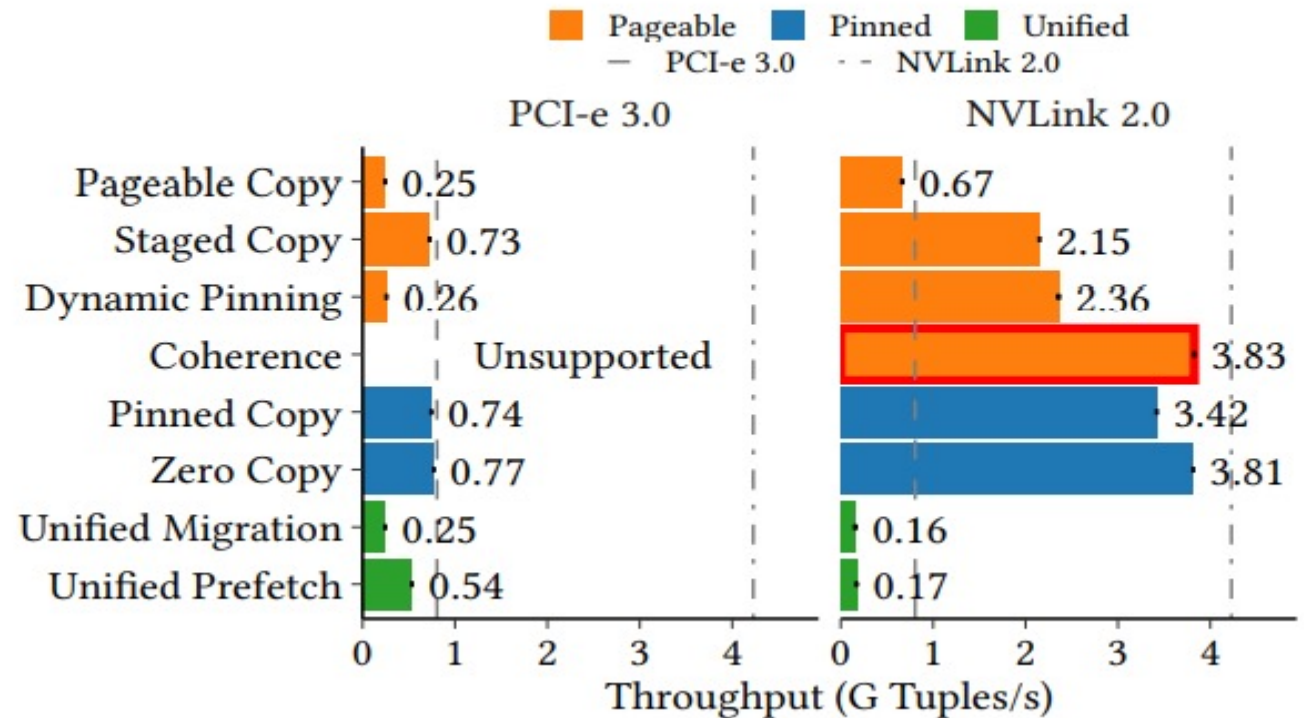
Table 2: Workload Overview.

| Property | A (from [10]) | B | C (from [54]) |
|--------------------|-----------------|-----------------|--------------------------|
| key / payload | 8 / 8 bytes | 8 / 8 bytes | 4 / 4 bytes |
| cardinality of R | 2^{27} tuples | 2^{18} tuples | $1024 \cdot 10^6$ tuples |
| cardinality of S | 2^{31} tuples | 2^{31} tuples | $1024 \cdot 10^6$ tuples |
| total size of R | 2 GiB | 4 MiB | 7.6 GiB |
| total size of S | 32 GiB | 32 GiB | 7.6 GiB |

EXPERIMENT RESULT (NVLINK VS OTHERS)

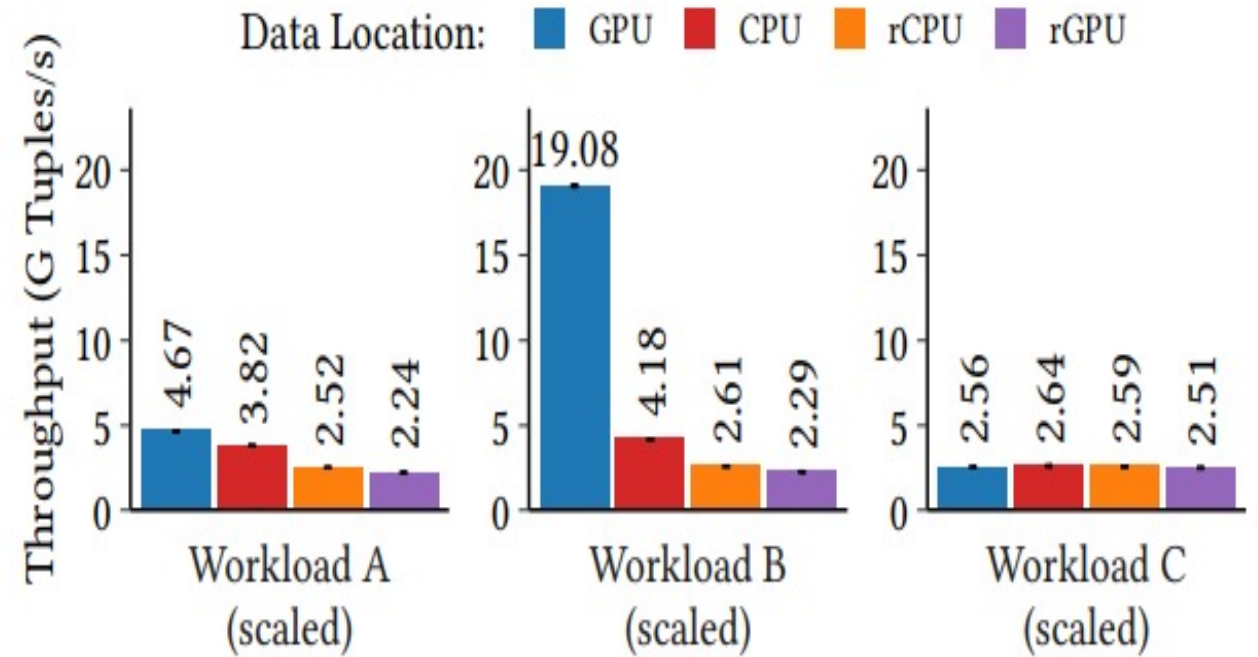
NVLink throughput is higher than PCI-e 3.0

Coherence produces the highest throughput



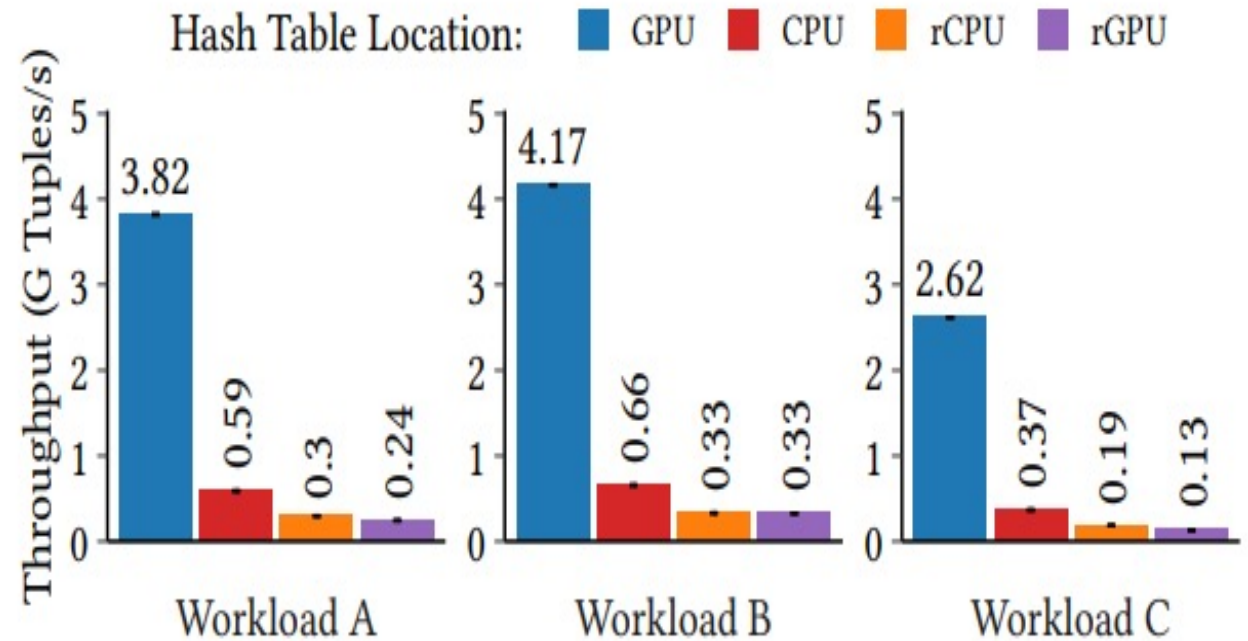
EXPERIMENT RESULT (DATA LOCATION)

Performance best when
data in 1 GPU memory



EXPERIMENT RESULT (HASHTABLE LOCATION)

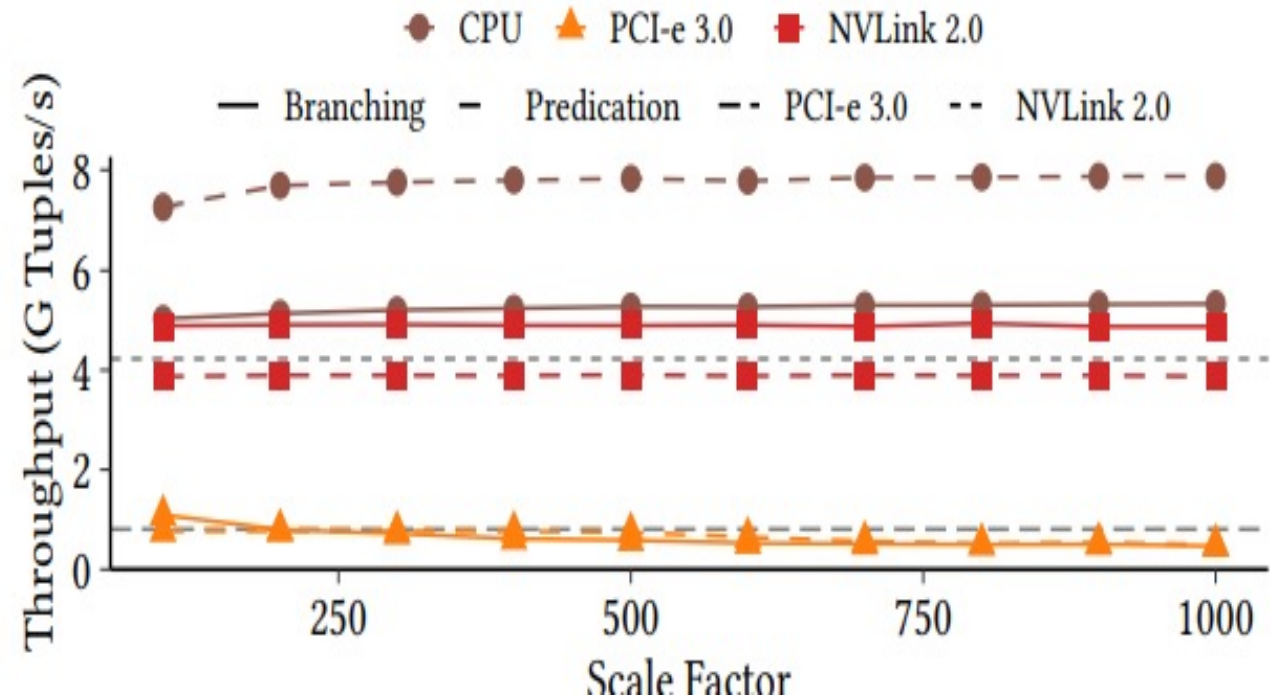
Performance best when
hash table in 1 GPU
memory



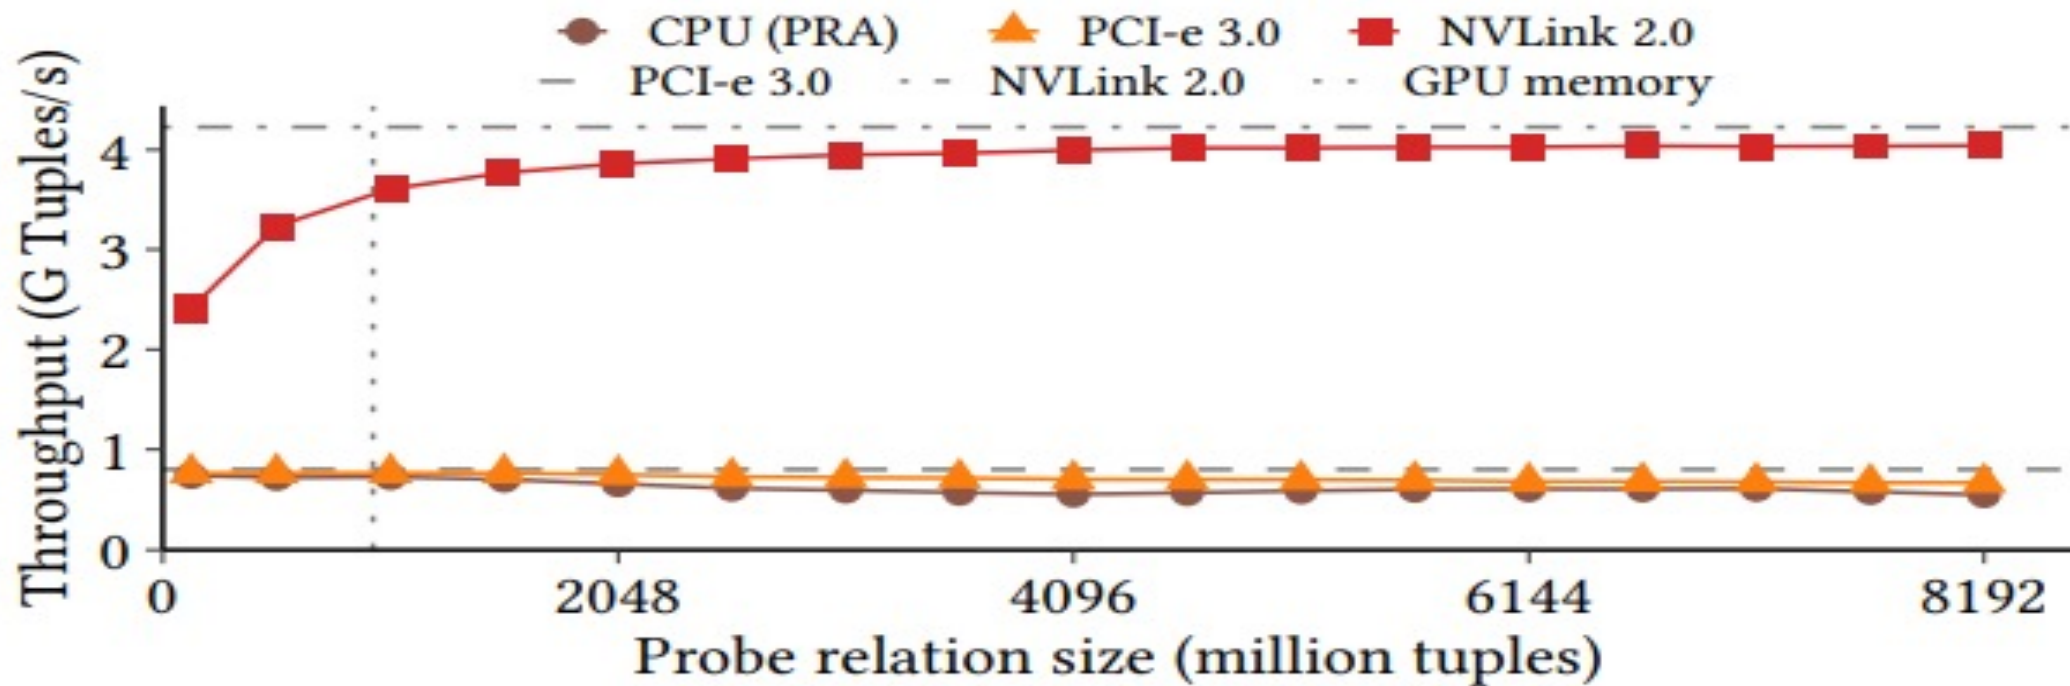
EXPERIMENT RESULT (SCALING DATA SIZE)

Interconnects. The CPU achieves the highest throughput, and outperforms NVLink 2.0

Branching vs. Predication. Branching performs better than predication on the GPU with NVLink 2.0.



EXPERIMENT RESULT (SCALING PROBE SIZE)



The throughput of NVLink 2.0 is the fastest

EXPERIMENT RESULT (SCALING BUILD SIZE)

NV Link provides best throughput

NVLink 2.0 with Hybrid Hash Table degrades gracefully

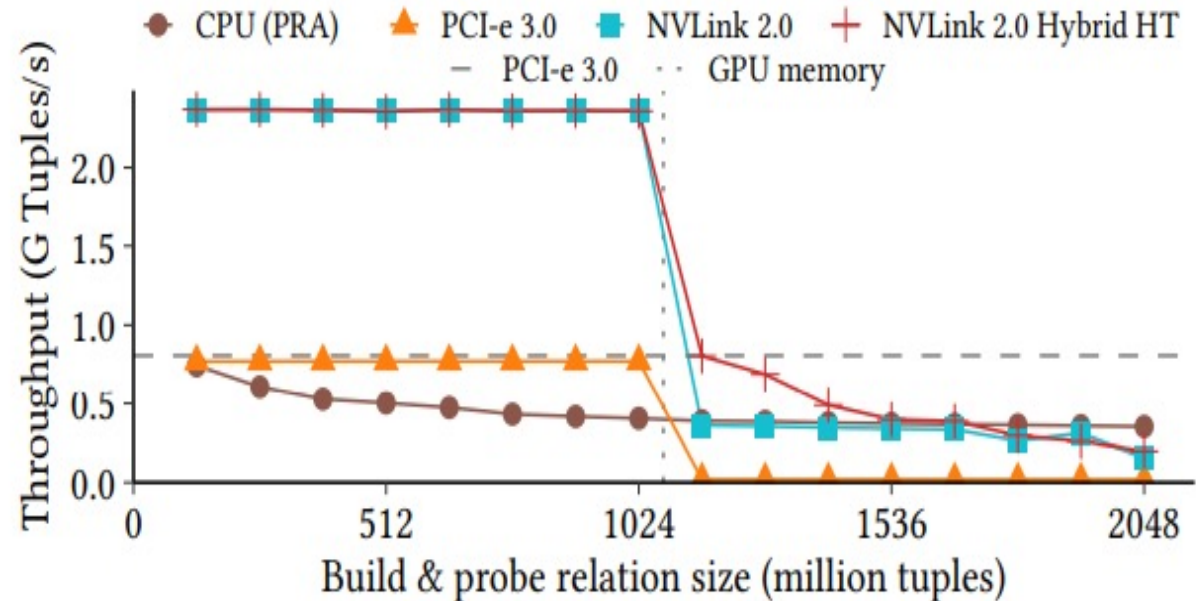


Figure 17: Scaling the build-side relation.

EXPERIMENT RESULT (BUILD TO PROBE RATIO)

The build phase takes 71% of the time

For larger ratios, the build-side takes up a smaller proportion of time

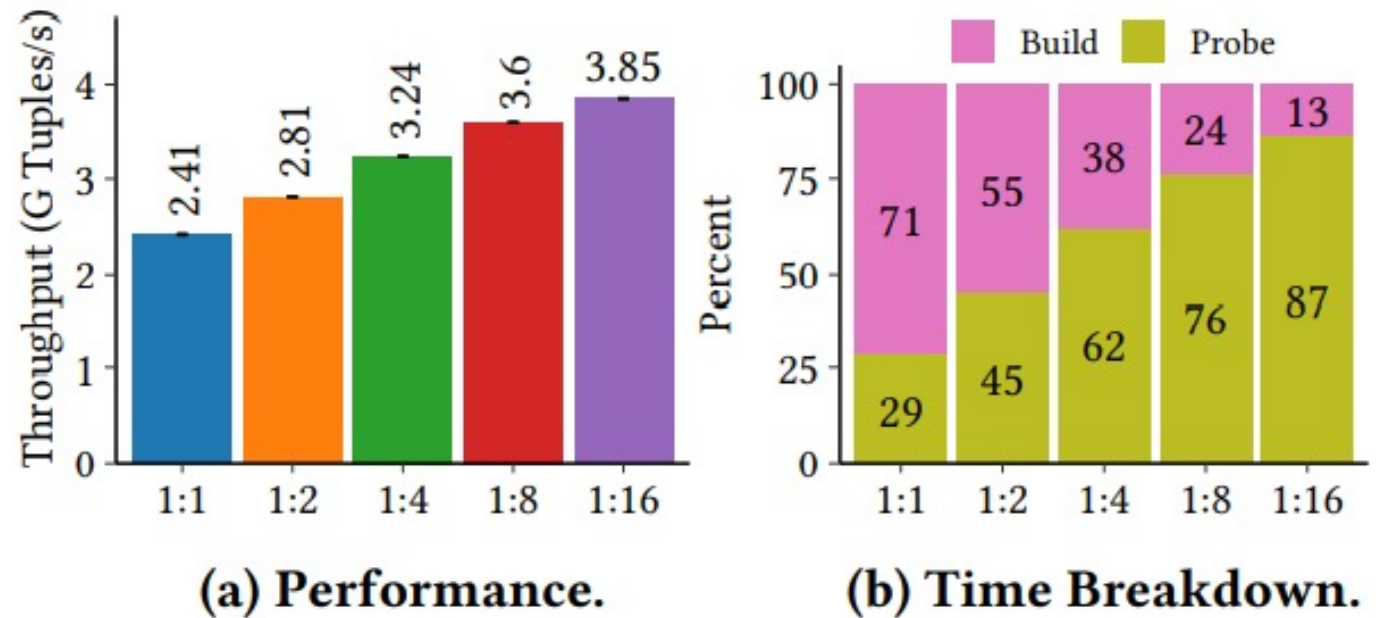


Figure 18: Different build-to-probe ratios on NVLink.

EXPERIMENT RESULT (BUILD DATA SKEW)

Higher skew
leads to a higher
throughput

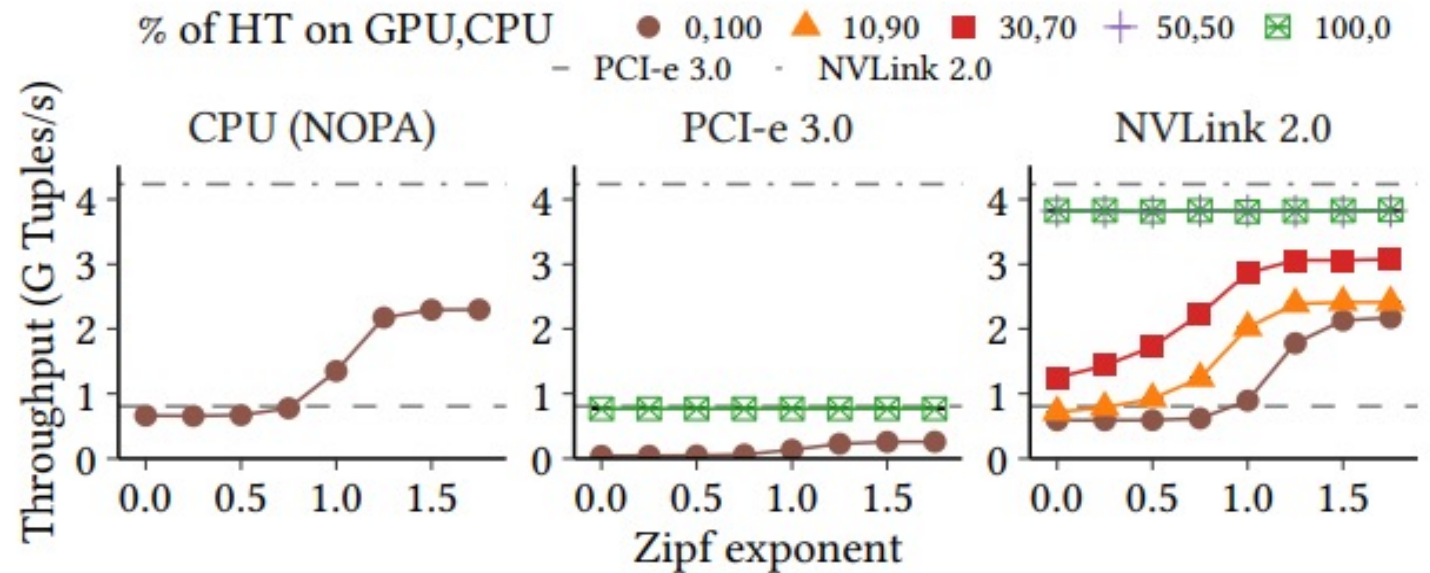


Figure 19: Join performance when the probe relation follows a Zipf distribution.

EXPERIMENT RESULT (JOIN SELECTIVITY)

Join throughput decreases with higher selectivity

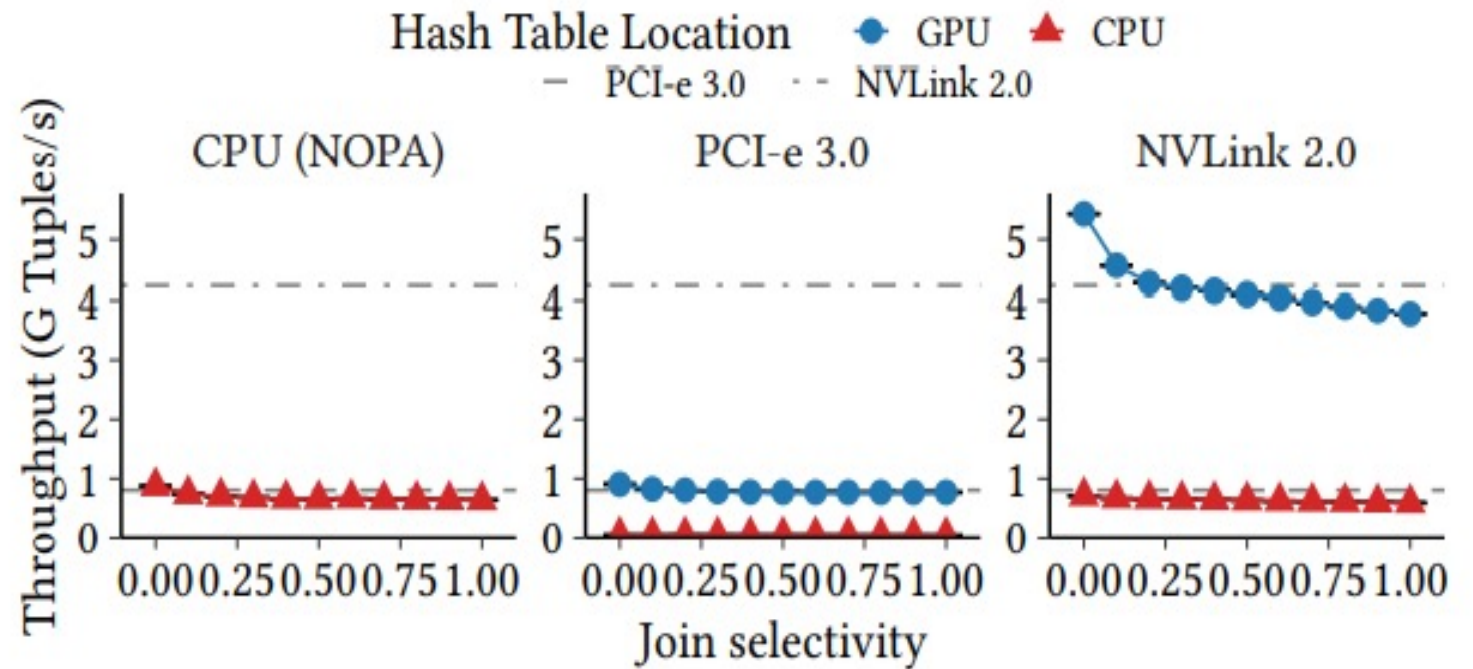
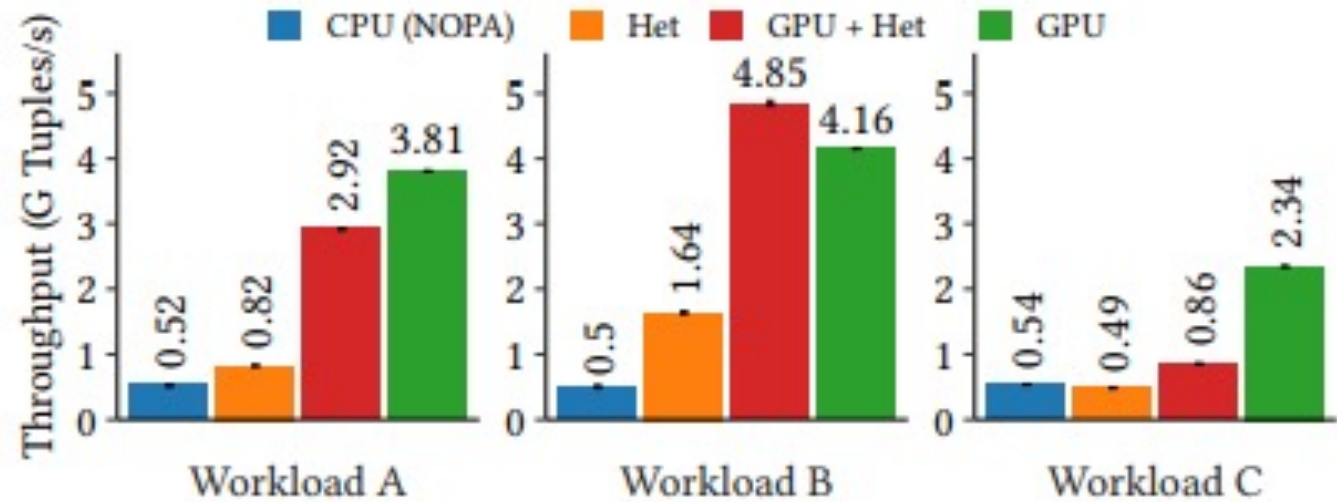


Figure 20: The effect of join selectivity on throughput.

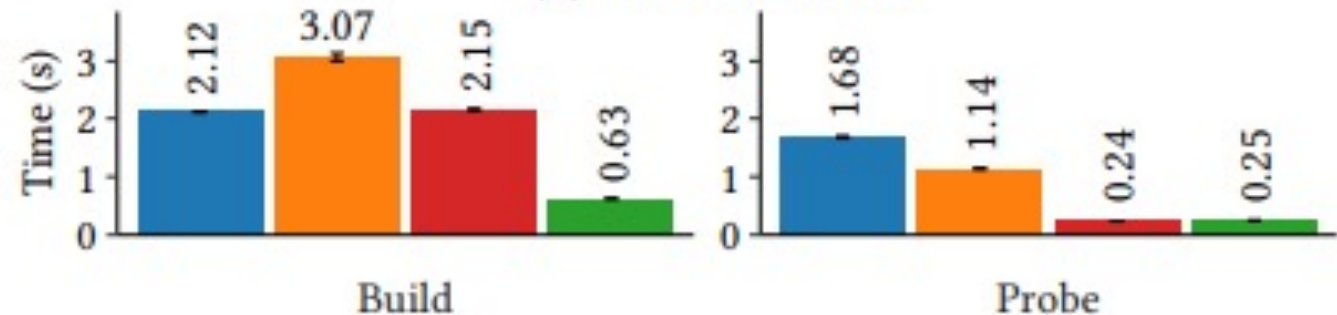
EXPERIMENT RESULT (CPU GPU CO PROCESSING SCALE UP)

1. Using a GPU always achieves the same or better throughput than the CPU-only strategy, and never decreases throughput.

2. GPU-only strategy achieves the best throughput for most of our workloads.



(a) Performance.



(b) Time per join phase in workload C (scaled).

INSIGHTS

- GPUs have high-bandwidth access to CPU memory
- GPUs can efficiently process large, out-of-core data
- GPUs are able to operate on out-of-core data structures, but should use GPU memory if possible
- Scaling-up co-processors with CPU + GPU makes performance more robust.
- Due to cache-coherence, memory pinning is no longer necessary to achieve high transfer bandwidth.
- Fair performance comparisons between GPUs vs. CPUs have become practical.

CONCLUSION

With fast interconnects, GPU acceleration becomes an attractive scale-up alternative that promises large speedups for databases.



THANK YOU